


Lecture 16 (Data Structures 2)

ADTs, BSTs

CS61B, Fall 2024 @ UC Berkeley

Slides credit: Josh Hug



Abstract Data Types

Lecture 16, CS61B, Fall 2024

Abstract Data Types

Binary Search Trees

- Derivation
- Definition
- contains
- Insert
- Hibbard deletion


Sets and Maps (are the same thing)

BST Implementation Tips

Interfaces vs. Implementation


In class:

- Developed ALists and SLLists.
- Created an interface List61B.
 - Modified AList and SLList to implement List61B.
 - List61B provided default methods.



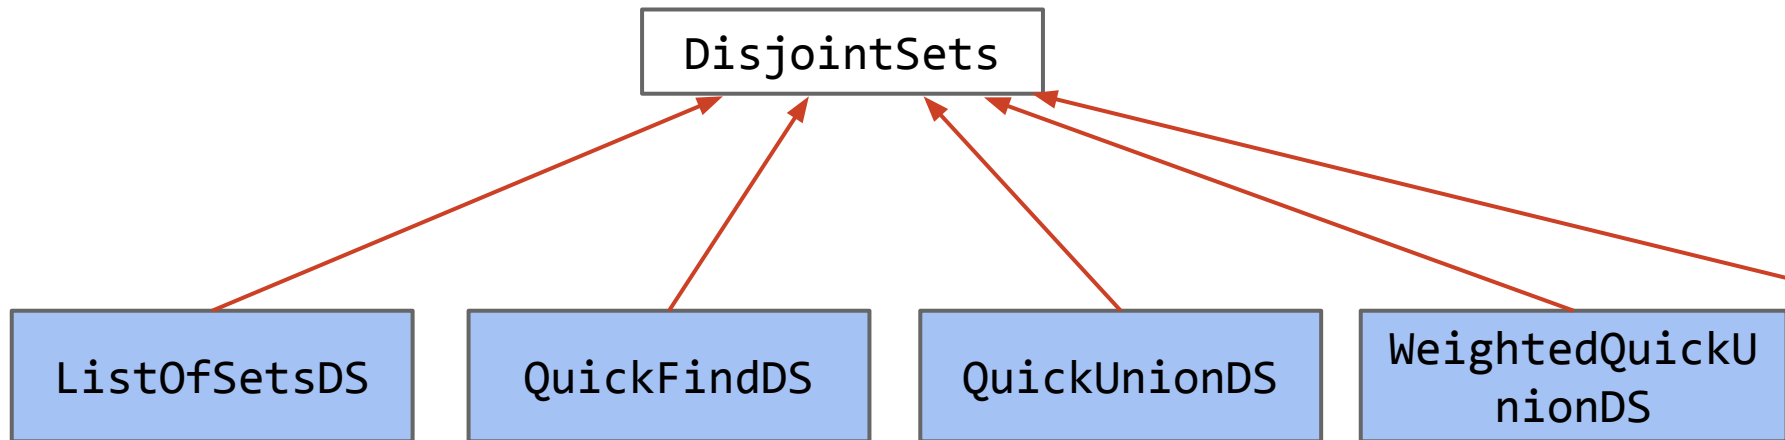
In projects:

- Developed ArrayDeque and LinkedListDeque.
 - Each class implemented the Deque interface.



Interfaces vs. Implementation

With `DisjointSets`, we saw a much richer set of possible implementations.




Abstract Data Types

An **Abstract Data Type (ADT)** is defined only by its operations, not by its implementation.

Deque ADT:

- `addFirst(Item x);`
- `addLast(Item x);`
- `boolean isEmpty();`
- `int size();`
- `printDeque();`
- `Item removeFirst();`
- `Item removeLast();`
- `Item get(int index);`




ArrayDeque and LinkedList Deque are implementations of the Deque ADT.



Another example of an ADT: The Stack

Recall, the Stack ADT supports the following operations:

- `push(int x)`: Puts `x` on top of the stack.
- `int pop()`: Removes and returns the top item from the stack.




Recall, the Stack ADT supports the following operations:

- `push(int x)`: Puts `x` on top of the stack.
- `int pop()`: Removes and returns the top item from the stack.

Which implementation do you think would result in faster overall performance?

- A. Linked List
- B. Array




The Stack ADT supports the following operations:

- `push(int x)`: Puts `x` on top of the stack.
- `int pop()`: Removes and returns the top item from the stack

Which implementation do you think would result in faster overall performance?

A. Linked List

B. Array



4

Both are about the same. No resizing for linked lists, so probably a lil faster.

The GrabBag ADT supports the following operations:

- `insert(int x)`: Inserts `x` into the grab bag.
- `int remove()`: Removes a random item from the bag.
- `int sample()`: Samples a random item from the bag (without removing!)
- `int size()`: Number of items in the bag.

Which implementation do you think would result in faster overall performance?

- A. Linked List
- B. Array

```
— insert(int x)
— remove()
— sample()
— size(int i)
```

The GrabBag ADT supports the following operations:

- `insert(int x)`: Inserts `x` into the grab bag.
- `int remove()`: Removes a random item from the bag.
- `int sample()`: Samples a random item from the bag (without removing!)
- `int size()`: Number of items in the bag.

Which implementation do you think would result in faster overall performance?

A. Linked List


B. Array

```
— insert(int x)
— remove()
— sample()
— size(int i)
```

One thing I particularly like about Java is the syntax differentiation between abstract data types and implementations.


- Note: Interfaces in Java aren't purely abstract as they can contain some implementation details, e.g. default methods.

Example: `List<Integer> L = new ArrayList<>();`



Among the most important interfaces in the `java.util` library are those that extend the `Collection` interface (btw interfaces can extend other interfaces).

- Lists of things.
- Sets of things.
- Mappings between items, e.g. jhug's grade is 88.4, or Creature c's north neighbor is a Plip.
 - Maps also known as associative arrays, associative lists (in Lisp), symbol tables, dictionaries (in Python).




Map Example

Maps are very handy tools for all sorts of tasks. Example: Counting words.

```
Map<String, Integer> m = new TreeMap<>();
String[] text = {"sumomo", "mo", "momo", "mo",
                "momo", "no", "uchi"};
for (String s : text) {
    int currentCount = m.getDefault(s, 0);
    m.put(s, currentCount + 1);
}
```

```
m = {}
text = ["sumomo", "mo", "momo", "mo", \
        "momo", "no", "uchi"]
for s in text:
    current_count = m.get(s, 0)
    m[s] = current_count + 1
```

Python
equivalent




sumomo	1
mo	2
momo	2
no	1
uchi	1

The built-in `java.util` package provides a number of useful:

- Interfaces: ADTs (lists, sets, maps, priority queues, etc.) and other stuff.
- Implementations: Concrete classes you can use.

Today, we'll learn the basic ideas behind the `TreeSet` and `TreeMap`.



Binary Search Trees: Derivation

Lecture 16, CS61B, Fall 2024

Abstract Data Types

Binary Search Trees


- **Derivation**
- Definition
- contains
- Insert
- Hibbard deletion

Sets and Maps (are the same thing)

BST Implementation Tips

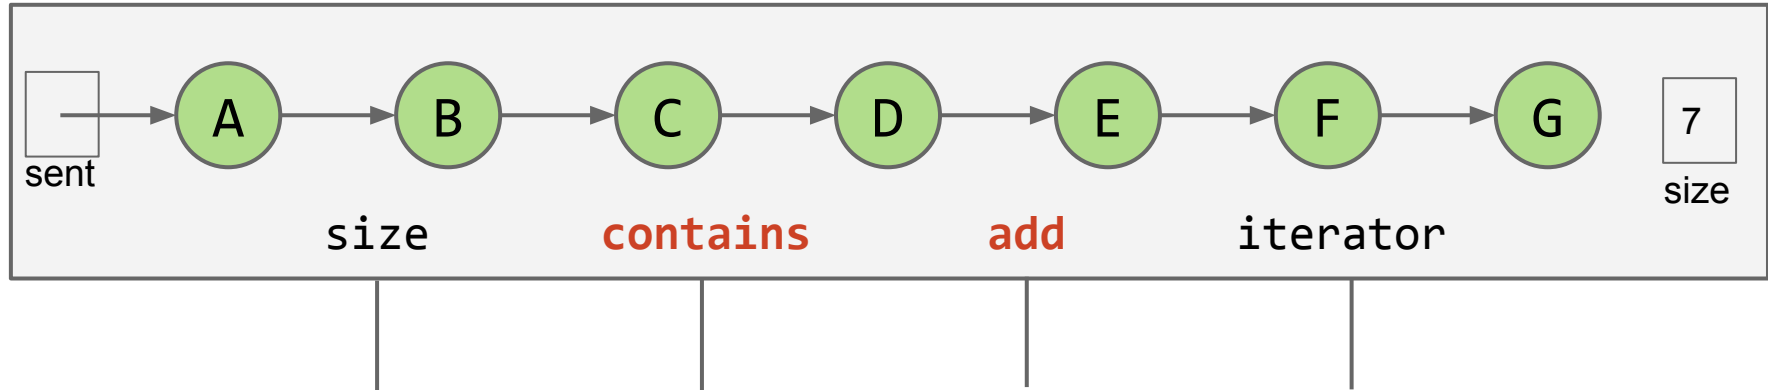
Analysis of an OrderedLinkedListSet<Character>

In an earlier lecture, we implemented a set based on [unordered arrays](#). For the **order linked list** set implementation below, name an operation that takes worst case linear time, i.e. $\Theta(N)$.




Analysis of an OrderedLinkedListSet<Character>

In an earlier lecture, we implemented a set based on [unordered arrays](#). For the **order linked list** set implementation below, name an operation that takes worst case linear time, i.e. $\Theta(N)$.



Fundamental Problem: Slow search, even though it's in order.


- Add (random) express lanes. [Skip List](#) (won't discuss in 61B)



Optimization: Change the Entry Point

Fundamental Problem: Slow search, even though it's in order.


- Move pointer to middle.



Optimization: Change the Entry Point, Flip Links

Fundamental Problem: Slow search, even though it's in order.


- Move pointer to middle and flip left links. Halved search time!



Optimization: Change the Entry Point, Flip Links

Fundamental Problem: Slow search, even though it's in order.


- How do we do even better?
- Dream big!



Optimization: Change Entry Point, Flip Links, Allow Big Jumps

Fundamental Problem: Slow search, even though it's in order.

- How do we do better?



Binary Search Trees: Definition

Lecture 16, CS61B, Fall 2024

Abstract Data Types

Binary Search Trees

- Derivation
- **Definition**
- contains
- Insert
- Hibbard deletion


Sets and Maps (are the same thing)

BST Implementation Tips

A tree consists of:

- A set of nodes.
- A set of edges that connect those nodes.
 - Constraint: There is exactly one path between any two nodes.

Green structures below are trees. Pink ones are not.



Rooted Trees and Rooted Binary Trees

In a rooted tree, we call one node the root.


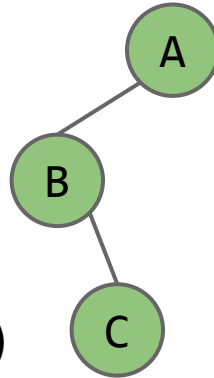
- Every node N except the root has exactly one parent, defined as the first node on the path from N to the root.
- Unlike [\(most\) real trees](#), the root is usually depicted at the top of the tree.
- A node with no child is called a leaf.

In a rooted binary tree, every node has either 0, 1, or 2 children (subtrees).



For each of these:

- A is the root.
- B is a child of A. (and C of B)
- A is a parent of B. (and B of C)



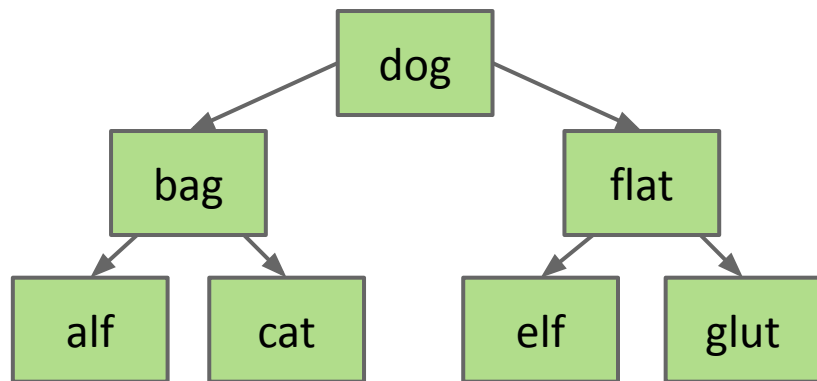
Not binary!

Binary Search Trees


A binary search tree is a rooted binary tree with the BST property.

BST Property. For every node X in the tree:

- Every key in the **left** subtree is **less** than X's key.
- Every key in the **right** subtree is **greater** than X's key.



Binary Search Tree



Binary Tree, but not a Binary Search Tree

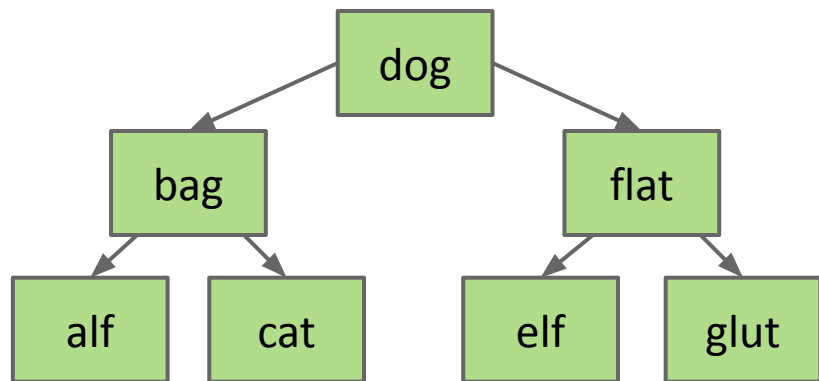
Binary Search Trees

Ordering must be complete, transitive, and antisymmetric. Given keys p and q :


- Exactly one of $p < q$ and $q < p$ are true.
- $p < q$ and $q < r$ imply $p < r$.

One consequence of these rules: No duplicate keys allowed!

- Keeps things simple. Most real world implementations follow this rule.



Binary Search Tree



Binary Tree, but not a Binary Search Tree

contains

Lecture 16, CS61B, Fall 2024

Abstract Data Types

Binary Search Trees

- Derivation
- Definition
- **contains**
- Insert
- Hibbard deletion

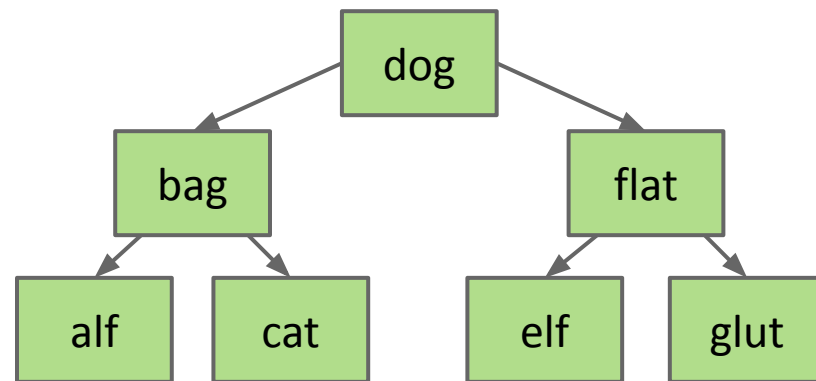
Sets and Maps (are the same thing)

BST Implementation Tips

Finding a searchKey in a BST (come back to this for the BST lab)

If searchKey equals T.key, return.

- If searchKey $<$ T.key, search T.left.
- If searchKey $>$ T.key, search T.right.

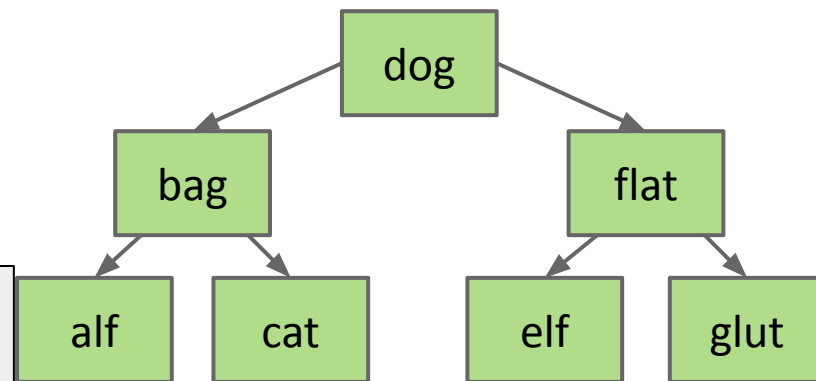


Finding a searchKey in a BST

If searchKey equals T.key, return.


- If searchKey < T.key, search T.left.
- If searchKey > T.key, search T.right.

```
static BST find(BST T, Key sk) {  
    if (T == null)  
        return null;  
    if (sk.equals(T.key))  
        return T;  
    else if (sk < T.key)  
        return find(T.left, sk);  
    else  
        return find(T.right, sk);  
}
```



What is the runtime to complete a search on a “bushy” BST in the worst case, where N is the number of nodes.


- A. $\Theta(\log N)$
- B. $\Theta(N)$
- C. $\Theta(N \log N)$
- D. $\Theta(N^2)$
- E. $\Theta(2^N)$



What is the runtime to complete a search on a “bushy” BST in the worst case, where N is the number of nodes.

A. $\Theta(\log N)$

Height of the tree is $\sim \log_2(N)$



Bushy BSTs are extremely fast.

- At 1 microsecond per operation, can find something from a tree of size 10^{300000} in one second.

Much (perhaps most?) computation is dedicated towards finding things in response to queries.

- It's a good thing that we can do such queries almost for free.

insert

Lecture 16, CS61B, Fall 2024

Abstract Data Types

Binary Search Trees

- Derivation
- Definition
- contains
- **insert**
- Hibbard deletion

Sets and Maps (are the same thing)

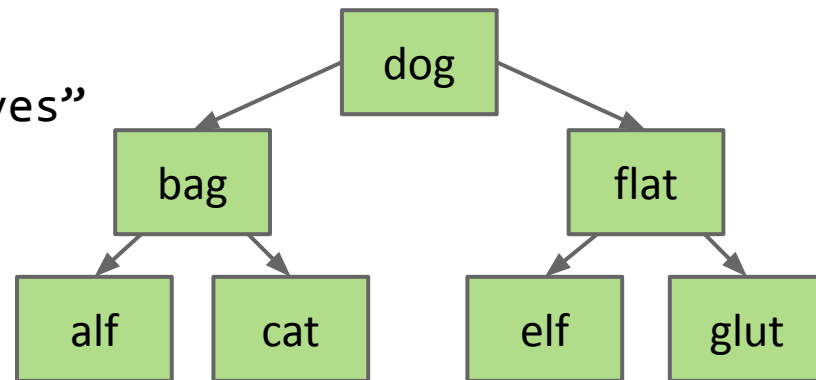
BST Implementation Tips

Inserting a New Key into a BST

Search for key.

- If found, do nothing.
- If not found:
 - Create new node.
 - Set appropriate link.

Example:
insert “eyes”

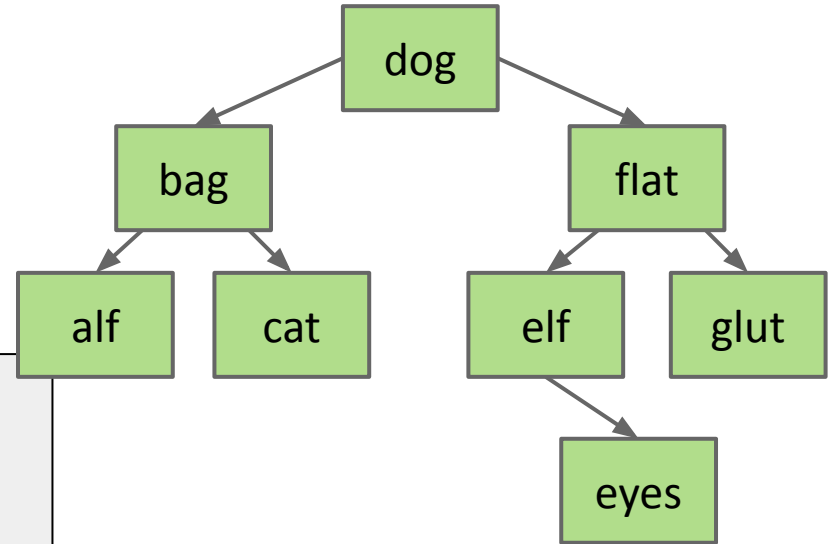


Inserting a New Key into a BST

Search for key.

- If found, do nothing.
- If not found:
 - Create new node.
 - Set appropriate link.

```
static BST insert(BST T, Key ik) {  
    if (T == null)  
        return new BST(ik);  
    if (ik < T.key)  
        T.left = insert(T.left, ik);  
    else if (ik > T.key)  
        T.right = insert(T.right, ik);  
    return T;  
}
```



Arms length recursion: A common rookie bad habit to avoid:

```
if (T.left == null)  
    T.left = new BST(ik);  
else if (T.right == null)  
    T.right = new BST(ik);
```

Avoid Arms-Length Recursion

```
if (T.left.left == null)
    T.left.left = new BST(ik);
else if (T.left.right == null)
    T.left.right = new BST(ik);
else if (T.right.left == null)
    T.right.left = new BST(ik);
else if (T.right.right == null)
    T.right.right = new BST(ik);
```

This base case is too complicated.
The recursion can take us further.

```
if (T.left == null)
    T.left = new BST(ik);
else if (T.right == null)
    T.right = new BST(ik);
```

Better, but still not the best base case.
Avoid arms-length recursion!

```
if (T == null)
    return new BST(ik);
```

The best base case.

Hibbard deletion

Lecture 16, CS61B, Fall 2024

Abstract Data Types

Binary Search Trees

- Derivation
- Definition
- contains
- Insert
- **Hibbard deletion**

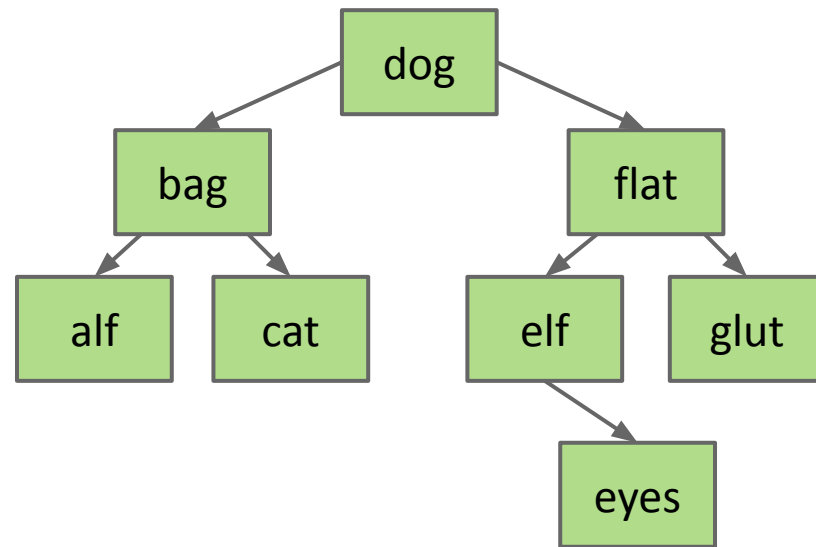
Sets and Maps (are the same thing)

BST Implementation Tips

Deleting from a BST

3 Cases:


- Deletion key has no children.
- Deletion key has one child.
- Deletion key has two children.



Case 1: Deleting from a BST: Key with no Children

Deletion key has no children (“glut”):

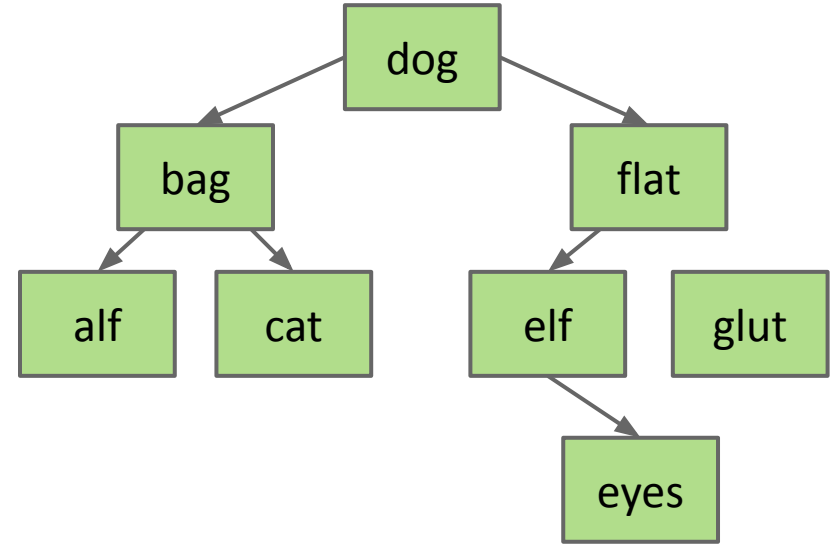
- Just sever the parent’s link.
- What happens to “glut” node?



Case 1: Deleting from a BST: Key with no Children

Deletion key has no children (“glut”):

- Just sever the parent’s link.
- What happens to “glut” node?
 - Garbage collected.

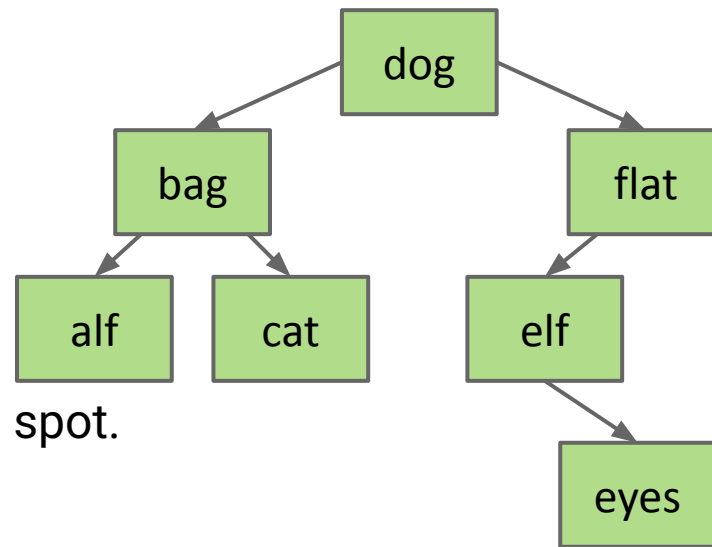


Case 2: Deleting from a BST: Key with one Child

Example: delete("flat"):

Goal:

- Maintain BST property.
- Flat's child definitely larger than dog.
 - Safe to just move that child into flat's spot.



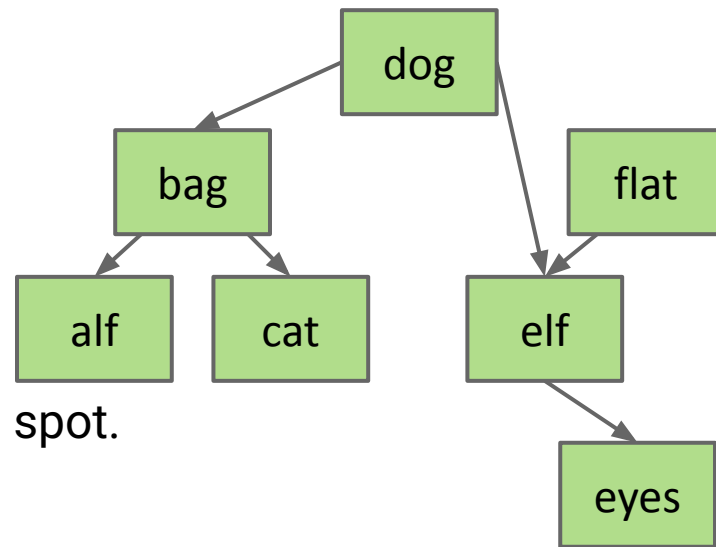
Thus: Move flat's parent's pointer to flat's child.

Case 2: Deleting from a BST: Key with one Child

Example: delete("flat"):

Goal:


- Maintain BST property.
- Flat's child definitely larger than dog.
 - Safe to just move that child into flat's spot.



Thus: Move flat's parent's pointer to flat's child.

- Flat will be garbage collected (along with its instance variables).

Delete k.




Case 3: Deleting from a BST: Deletion with two Children (Hibbard)

Example: delete("dog")

Goal:

- Find a new root node.
- Must be $>$ than everything in left subtree.
- Must be $<$ than everything right subtree.

Would bag work?

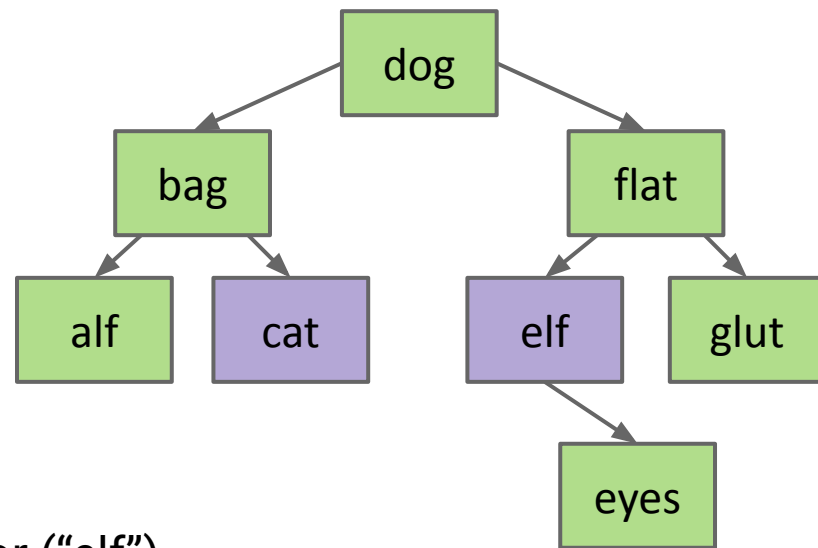


Case 3: Deleting from a BST: Deletion with two Children (Hibbard)

Example: delete("dog")

Goal:

- Find a new root node.
- Must be $>$ than everything in left subtree.
- Must be $<$ than everything right subtree.




Choose either predecessor ("cat") or successor ("elf").

- Delete "cat" or "elf", and stick new copy in the root position:
 - This deletion guaranteed to be either case 1 or 2. Why?
- This strategy is sometimes known as "Hibbard deletion".

Hard Challenge (Hopefully Now Easy)


Delete k.



Hard Challenge (Hopefully Now Easy)

Delete k. Two solutions: Either promote g or m to be in the root.


- Below, solution for g is shown.



Hard Challenge (Hopefully Now Easy)

Two solutions: Either promote g or m to be in the root.

- Below, solution for g is shown.



Sets and Maps (are the same thing)

Lecture 16, CS61B, Fall 2024

Abstract Data Types

Binary Search Trees

- Derivation
- Definition
- contains
- Insert
- Hibbard deletion



Sets and Maps (are the same thing)

BST Implementation Tips

Sets vs. Maps

Can think of the BST below as representing a Set:



- {mo, no, sumomo, uchi, momo}



Sets vs. Maps

Can think of the BST below as representing a Set:

- {mo, no, sumomo, uchi, momo}



sumomo
mo
momo
no
uchi



But what if we wanted to represent a mapping of word counts?

????

sumomo	1
mo	2
momo	2
no	1
uchi	1

Sets vs. Maps

To represent maps, just have each BST node store key/value pairs.



sumomo	1
mo	2
momo	2
no	1
uchi	1

Note: No efficient way to look up by value.

- Example: Cannot find all the keys with value = 1 without iterating over ALL nodes. This is fine.

Abstract data types (ADTs) are defined in terms of operations, not implementation.

Several useful ADTs: Disjoint Sets, Map, Set, List.

- Java provides Map, Set, List interfaces, along with several implementations.

We've seen two ways to implement a Set (or Map): ArraySet and using a BST.

- ArraySet: $\Theta(N)$ operations in the worst case.
- BST: $\Theta(\log N)$ operations in the worst case if tree is balanced.

BST Implementations:

- Search and insert are straightforward (but insert is a little tricky).
- Deletion is more challenging. Typical approach is "Hibbard deletion".

BST Implementation Tips

Lecture 16, CS61B, Fall 2024

Abstract Data Types

Binary Search Trees

- Derivation
- Definition
- contains
- Insert
- Hibbard deletion

Sets and Maps (are the same thing)

BST Implementation Tips

Tips for BST Lab

- Code from class was “naked recursion”. Your BSTMap will not be.
- For each public method, e.g. `put(K key, V value)`, create a private recursive method, e.g. `put(K key, V value, Node n)`
- When inserting, always set left/right pointers, even if nothing is actually changing.
- Avoid “arms length base cases”. Don’t check if left or right is null!

```
static BST insert(BST T, Key ik) {  
    if (T == null)  
        return new BST(ik);  
    if (ik < T.label())  
        T.left = insert(T.left, ik);  
    else if (ik > T.label())  
        T.right = insert(T.right, ik);  
    return T;  
}
```

Always set, even if
nothing changes!

Avoid “arms length base cases”.

```
if (T.left == null)  
    T.left = new BST(ik);  
else if (T.right == null)  
    T.right = new BST(ik);
```